

**CO2 – AT3**  
**MLA 0201 – Fundamentals of Machine Learning**

**Annexure A**

**Debugging & Optimisation Task**

Q1. A real estate company developed a linear regression model to predict house prices, but the model is showing high prediction error.

- (i) Identify possible causes for the high error (e.g., multicollinearity, outliers, poor feature selection).
- (ii) Analyze the dataset/model behavior to detect these issues.
- (iii) Apply suitable optimization techniques (such as feature scaling, removing outliers, or adding relevant features) to improve the model performance. (5 Marks-8 Mins)

Q2. An email classification system using a linear model is incorrectly classifying several spam emails as non-spam.

- (i) Identify the reasons for misclassification (e.g., improper decision boundary, imbalanced data).
- (ii) Examine the current decision boundary and classification errors.
- (iii) Refine the model using optimization techniques such as adjusting thresholds, feature engineering, or re-training with balanced data. (5 Marks-8 Mins)

Q3. A Naïve Bayes classifier is giving poor accuracy due to missing or zero probability values in the dataset.

- (i) Identify the issue related to missing or zero probabilities.
- (ii) Analyze how these values affect classification results.
- (iii) Apply appropriate debugging and optimization methods such as Laplace smoothing to improve accuracy. (5 Marks-8 Mins)

Q4. A discriminant function is unable to clearly separate two classes, leading to incorrect classifications.

- (i) Identify possible reasons for poor class separation (e.g., overlapping features, incorrect weights).
- (ii) Analyze the decision boundary and feature contribution.

(iii) Improve the discriminant function using optimization strategies such as adjusting weights, feature transformation, or normalization.(5 Marks-8 Mins)

Q5.A Bayesian logistic regression model developed for customer churn prediction is suffering from overfitting, resulting in poor generalization.

- (i) Identify signs and causes of overfitting in the model.
- (ii) Analyze model performance on training vs testing data.
- (iii) Apply debugging and regularization techniques (such as L1/L2 regularization or cross validation) to improve model performance.(5 Marks-8 Mins)

Q6.A linear regression model built for predicting sales is performing poorly because the features (e.g., advertising budget in thousands and number of outlets in units) are on different scales. (i) Identify how feature scaling issues affect model performance.

- (ii) Analyze the impact of unscaled features on regression coefficients.
- (iii) Apply appropriate optimization techniques such as normalization or standardization to improve the model.(5 Marks-8 Mins)

Q7.A binary classification model for fraud detection is biased towards the majority class, leading to poor detection of fraudulent cases.

- (i) Identify the problem caused by class imbalance.
- (ii) Evaluate model performance using suitable metrics (precision, recall, F1-score).
- (iii) Apply optimization techniques such as resampling (oversampling/undersampling) or class weighting to improve performance.(5 Marks-8 Mins)

Q8.A Naïve Bayes classifier is underperforming because some input features are highly correlated, violating the independence assumption.

- (i) Identify how correlated features affect Naïve Bayes performance.
- (ii) Analyze feature relationships in the dataset.
- (iii) Apply optimization techniques such as feature selection or dimensionality reduction to improve accuracy. (5 Marks-8 Mins)

Q9.A logistic regression model is unable to clearly separate two classes due to overlapping feature distributions.

- (i) Identify the cause of overlapping decision boundaries.
- (ii) Analyze the feature space and classification errors.
- (iii) Improve the model using techniques such as feature transformation, polynomial features, or regularization. (5 Marks-8 Mins)

Q10.A discriminant analysis model is too simple and fails to capture the complexity of the

dataset, resulting in low accuracy.

- (i) Identify the signs of underfitting in the model.
- (ii) Analyze model bias and feature limitations.
- (iii) Apply optimization strategies such as adding features, increasing model complexity, or using nonlinear transformations. (5 Marks-8 Mins)

## **1. Linear Regression – High House Price Prediction**

### **Error (i) Problem Identification**

High prediction error may be caused by:

- **Multicollinearity** – two or more features are highly correlated.
- **Outliers** – unusually expensive or cheap houses can distort the regression line.
- **Poor feature selection** – important features such as location, number of rooms, or house age may be missing.
- **Different feature scales** – features with very different ranges can affect optimization.

**Non-linear relationships** – house price may not increase linearly with every feature. (ii)

### **Debugging and Analysis**

The dataset can be analyzed using:

- Correlation matrix to detect highly correlated features.
- Box plots or IQR to identify outliers.
- Train-test split to compare training and testing errors.
- MAE, MSE and RMSE to measure prediction performance.
- Residual analysis to check whether prediction errors are randomly distributed.

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
pred = model.predict(X_test)
```

```
print("RMSE:", mean_squared_error(y_test, pred) ** 0.5)
```

```
print("R2 Score:", r2_score(y_test, pred))
```

```
print("Correlation:\n", X.corr())
```

### **(iii) Optimization**

The model can be improved by:

- Removing extreme outliers using IQR.
- Removing one of highly correlated features.
- Adding relevant features such as location, house age and number of rooms. •

Applying **StandardScaler** or **MinMaxScaler**.

- Using regularization such as Ridge or Lasso.

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import Ridge
```

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
model = Ridge(alpha=1.0)
```

```
model.fit(X_train, y_train)
```

**Result:** Removing noise, selecting useful features and regularizing the model reduces prediction error and improves generalization.

## **2. Linear Model – Spam Email Misclassification**

### **(i) Problem Identification**

Spam emails may be classified as non-spam because of:

- Incorrect decision boundary.
- Insufficient or irrelevant features.

- Imbalanced spam/non-spam data.
- Similar vocabulary between spam and legitimate emails.
- Poorly selected classification threshold.

## **(ii) Debugging and Analysis**

Analyze the **confusion matrix**:

- **True Positive (TP)**: Spam correctly identified.
- **True Negative (TN)**: Non-spam correctly identified.
- **False Positive (FP)**: Normal email classified as spam.
- **False Negative (FN)**: Spam classified as non-spam.

Since missing spam is the main problem, **recall for the spam class** is

important. from sklearn.metrics import confusion\_matrix,

classification\_report

pred = model.predict(X\_test)

print(confusion\_matrix(y\_test, pred))

print(classification\_report(y\_test, pred))

## **(iii) Optimization**

- Adjust the classification threshold.
- Add features such as spam keywords, URL count, special-character count and email length.
- Balance the training dataset.
- Use class weights.
- Retrain the model using better features.

```
from sklearn.linear_model import LogisticRegression

model = LogisticRegression(class_weight="balanced")

model.fit(X_train, y_train)
```

**Result:** Better features and balanced training reduce false negatives and improve spam detection recall.

### **3. Naïve Bayes – Missing or Zero Probabilities**

#### **(i) Problem Identification**

Naïve Bayes calculates:

If any conditional probability becomes **zero**, the entire product becomes zero.

Missing values can also produce incorrect probability calculations.

#### **(ii) Debugging and Analysis**

Check:

- Missing values in the dataset.
- Features having zero occurrences for a particular class.
- Probability tables generated by the classifier.
- Classification results before and after smoothing.

For example:

Then:

even if other words strongly indicate spam.

#### **(iii) Optimization – Laplace Smoothing**

Laplace smoothing prevents zero probabilities:

where  $V$  is the number of possible feature values.

```
from sklearn.naive_bayes import MultinomialNB

model = MultinomialNB(alpha=1.0)

model.fit(X_train, y_train)

pred = model.predict(X_test)
```

Here,  $\alpha=1$  provides Laplace smoothing.

**Result:** Zero probabilities are avoided, making the classifier more robust and improving classification accuracy.

## **4. Discriminant Function – Poor Class Separation**

### **(i) Problem Identification**

Poor class separation can occur because:

- Features of the two classes overlap.
- Incorrect discriminant weights are used.
- Features have different scales.
- Irrelevant features introduce noise.
- The selected features do not sufficiently distinguish the classes.

A discriminant function can be represented as:

The class is selected based on the value of .

### **(ii) Debugging and Analysis**

Analyze:

- Feature distributions of both classes.
- Decision boundary.
- Contribution of each feature through its weight.
- Misclassified samples.

- Correlation between features.

A large absolute weight indicates that the corresponding feature has a strong influence on the decision.

### **(iii) Optimization**

- Normalize or standardize features.
- Adjust or retrain weights.
- Remove irrelevant features.
- Apply feature transformation.
- Add more discriminative features.

```
from sklearn.preprocessing import StandardScaler

from sklearn.discriminant_analysis import

LinearDiscriminantAnalysis scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)

model = LinearDiscriminantAnalysis()

model.fit(X_train, y_train)
```

**Result:** Better-scaled and more informative features produce a clearer decision boundary and improve classification accuracy.

## **5. Bayesian Logistic Regression – Customer Churn**

### **Overfitting (i) Problem Identification**

Signs of overfitting:

- Very high training accuracy.
- Much lower testing accuracy.

- Large difference between training and validation performance. •

Model performs poorly on unseen customers.

Causes include:

- Too many features.
- No regularization.
- Small training dataset.
- Noisy or irrelevant features.
- Excessive model complexity.

## **(ii) Debugging and Analysis**

Compare:

```
train_score = model.score(X_train, y_train)
```

```
test_score = model.score(X_test, y_test)
```

```
print("Training:", train_score)
```

```
print("Testing:", test_score)
```

For example:

**Training accuracy = 98%**

**Testing accuracy = 72%**

This large difference indicates overfitting.

## **(iii) Optimization**

Use:

- **L1 regularization** for feature selection.
- **L2 regularization** to reduce large coefficients.
- Cross-validation.

- Feature selection.
- More training data if available.

```
from sklearn.linear_model import LogisticRegression

model = LogisticRegression(

    penalty="l2",

    C=0.1,

    max_iter=1000

)

model.fit(X_train, y_train)
```

Cross-validation can be used to select the best regularization strength.

**Result:** Regularization controls model complexity and improves performance on unseen customer data.

## 6. Linear Regression – Features on Different Scales

### (i) Problem Identification

Suppose the features are:

- Advertising budget = ₹10,000–₹1,00,000
- Number of outlets = 1–100

These features have very different scales.

Large-scale features can dominate the optimization process and cause inefficient gradient-based learning.

### (ii) Debugging and Analysis

Before scaling:

The coefficients may have very different numerical magnitudes because the input features use different units.

Check:

```
print(X_train.min())
```

```
print(X_train.max())
```

If the ranges differ significantly, scaling is appropriate.

### **(iii) Optimization**

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import LinearRegression
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
model = LinearRegression()
```

```
model.fit(X_train_scaled, y_train)
```

**Result:** Scaling puts features on comparable ranges, improves numerical stability and makes optimization more efficient.

---

## **7. Fraud Detection – Class Imbalance**

### **(i) Problem Identification**

Fraud datasets usually contain many legitimate transactions and very few fraudulent transactions.

Example:

- Normal = 9,500

- Fraud = 500

A model predicting every transaction as normal can achieve **95% accuracy**, but it completely fails to detect fraud.

Therefore, accuracy alone is misleading.

## **(ii) Debugging and Analysis**

Use:

**Precision**

**Recall**

**F1-score**

```
from sklearn.metrics import classification_report  
  
print(classification_report(y_test, pred))
```

For fraud detection, **recall is especially important** because missing fraudulent transactions is costly.

## **(iii) Optimization**

Possible solutions:

- Oversampling the minority class.
- Undersampling the majority class.
- SMOTE.
- Class weighting.

```
from sklearn.linear_model import LogisticRegression  
  
model = LogisticRegression(class_weight="balanced")  
  
model.fit(X_train, y_train)
```

**Result:** The model gives greater importance to fraudulent transactions and improves fraud detection recall and F1-score.

## 8. Naïve Bayes – Highly Correlated Features

### (i) Problem Identification

Naïve Bayes assumes that features are conditionally independent:

When features are highly correlated, this assumption is violated.

For example:

- annual\_income
- monthly\_income

contain related information.

Naïve Bayes may effectively count the same information multiple times. **(ii) Debugging and Analysis**

Use a correlation matrix:

```
import pandas as pd  
  
correlation = X.corr()  
  
print(correlation)
```

Features with correlation close to **+1 or -1** may contain redundant information. **(iii) Optimization**

Use:

- Feature selection.
- Remove highly correlated features.

- PCA for dimensionality reduction.

- Select the most informative features.

```
from sklearn.feature_selection import SelectKBest, chi2
```

```
selector = SelectKBest(chi2, k=10)
```

```
X_train_new = selector.fit_transform(X_train, y_train)
```

```
X_test_new = selector.transform(X_test)
```

**Result:** Removing redundant features reduces violation of the independence assumption and can improve Naïve Bayes performance.

---

## 9. Logistic Regression – Overlapping Classes

### (i) Problem Identification

Logistic regression normally creates a linear decision boundary:

If the feature distributions of two classes overlap heavily, a straight line cannot separate them effectively.

### (ii) Debugging and Analysis

Analyze:

- Feature distributions.
- Misclassified samples.
- Confusion matrix.
- Decision boundary.

- Relationship between features and target.

A high number of both false positives and false negatives indicates poor class separation.

### **(iii) Optimization**

Use:

- Feature transformation.
- Polynomial features.
- Feature interaction.
- Regularization.
- Feature selection.

```
from sklearn.preprocessing import PolynomialFeatures
```

```
from sklearn.linear_model import LogisticRegression
```

```
poly = PolynomialFeatures(degree=2)
```

```
X_train_poly = poly.fit_transform(X_train)
```

```
X_test_poly = poly.transform(X_test)
```

```
model = LogisticRegression(C=1.0, max_iter=1000)
```

```
model.fit(X_train_poly, y_train)
```

Polynomial features allow logistic regression to model more complex boundaries.

**Result:** Feature transformation increases model flexibility and can improve classification when classes are not linearly separable.

## **10. Discriminant Analysis – Underfitting**

### **(i) Problem Identification**

Underfitting occurs when the model is **too simple** to represent the actual relationship in the data.

Signs include:

- Low training accuracy.
- Low testing accuracy.
- High bias.
- Important features are missing.
- Model cannot capture nonlinear relationships.

Example:

**Training accuracy = 65%**

**Testing accuracy = 63%**

Since both are low, the model is likely underfitting.

## **(ii) Debugging and Analysis**

Compare training and testing performance:

```
print("Training Accuracy:",model.score(X_train, y_train))
```

```
print("Testing Accuracy:",model.score(X_test, y_test))
```

If both scores are low, investigate:

- Number of features.
- Feature quality.
- Model complexity.
- Nonlinear relationships.
- Data representation.

## **(iii) Optimization**

Improve the model by:

- Adding relevant features.

- Applying nonlinear transformations.
- Creating interaction features.
- Increasing model complexity.
- Using a more flexible discriminant model if necessary.

```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly = PolynomialFeatures(degree=2)
```

```
X_train_new = poly.fit_transform(X_train)
```

```
X_test_new = poly.transform(X_test)
```

**Result:** Adding informative features and increasing model flexibility reduces bias and improves both training and testing performance.